

Machine Learning

Week 2

Linear Regression

ECE

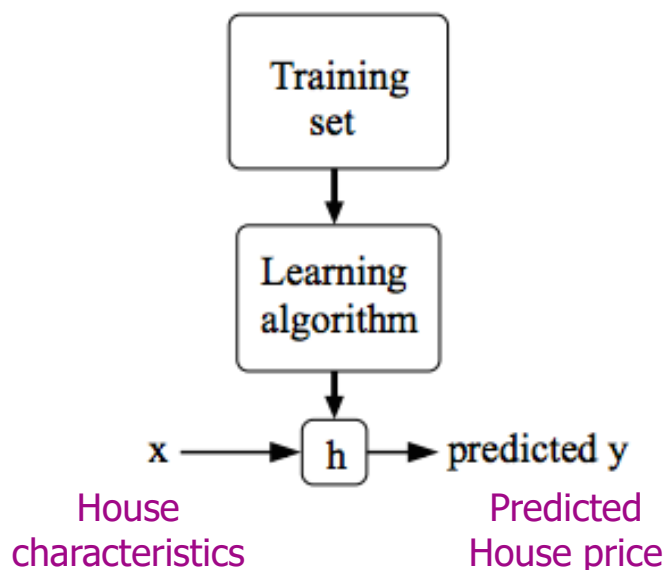
14/01/2019

Course overview

1. Introduction to Linear Regression
2. Simple Linear Regression
3. Multiple Linear Regression
4. Evaluation of a Linear Regression Model
5. Practical Work

Model Representation

- $x^{(i)}$ denotes the “input” variables (house characteristics)
- $y^{(i)}$ denotes the “output” or target variable that we are trying to predict (price)
- A pair $(x^{(i)}, y^{(i)})$ is called a training example
- A list of m training examples $(x^{(i)}, y^{(i)}); i=1, \dots, m$ —is called a training set




$$\mathbf{h}: \mathbf{X} \rightarrow \mathbf{Y}$$

Hypothesis or function that takes as input the house's characteristics to estimate its price.

1.2


Simple/Univariate Linear Regression

Simple Linear Regression: Model



$$y = \theta_0 + \theta_1 x + \varepsilon$$

Simple Linear Regression: Model Evaluation



Each line has different parameters $\theta = (\theta_0, \theta_1)$
→ different errors

- Which line is the best fit ?
- what should a good function $h_{\theta}(x)$ minimize?
 - Sum error on all data points
 - Sum abs(error) on all data points
 - Sum error² on all data points

Simple Linear Regression: Model Evaluation

- Sum error on all data points

Sum error = 0




- Sum abs(error) on all data points

Sum abs(error) > 0




- Sum error^2 on all data points

Sum error^2 > 0



Simple Linear Regression: Model Evaluation


- The Sum of Squared Errors (SSE) is also called Residual Sum of Squares (RSS)



$$SSE(\theta_0, \theta_1) =$$

$$\begin{aligned} & (y^{(1)} - (\theta_0 + \theta_1 * x^{(1)}))^2 + \\ & (y^{(2)} - (\theta_0 + \theta_1 * x^{(2)}))^2 + \\ & \dots\dots\dots + \\ & (y^{(m)} - (\theta_0 + \theta_1 * x^{(m)}))^2 \end{aligned}$$

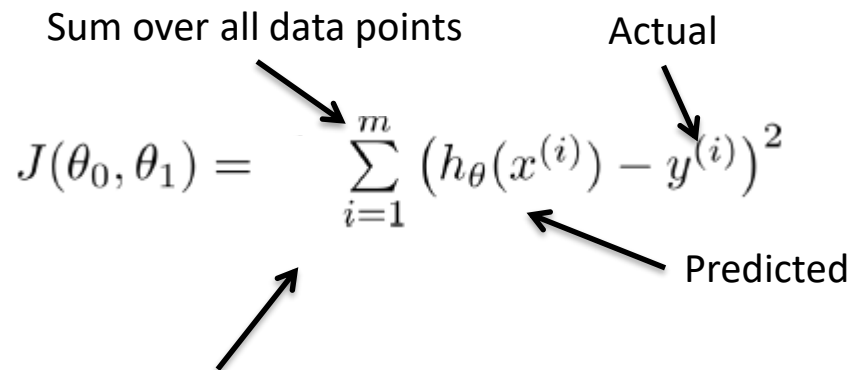
Simple Linear Regression: Model Evaluation



- The green line is a better fit
- Let's see how to find the best line automatically

Simple Linear Regression: Cost function

- The best hypothesis $h_{\theta}(x)$ is the one that minimizes the cost function


$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- $1/m$ - means we determine the average
- $1/2m$, the 2 makes the math a bit easier, and doesn't change the weights θ we determine at all (i.e. half the smallest value is still the smallest value!)
- The learning algorithm should find $\theta^* = (\theta_0^*, \theta_1^*)$ that minimizes this cost
- Several algorithms:
 - Gradient descent
 - Ordinary least square (OLS): Used in the linear regression in python (sklearn)


Gradient Descent Algorithm: Intuition

- Let's assume that we have one parameter θ , and that we want to minimize $J(\theta)$

$$h_{\theta}(x) = \theta x$$

- GD starts by a random initial θ and iteratively update it to get towards θ^* .

In this case, the derivative
(gradient) $\partial J(\theta) / \partial \theta < 0$.
 $\theta^A - \alpha * \partial J(\theta) / \partial \theta > \theta^A$
 θ^A is moving to the right
 θ is increasing




In this case, the derivative
(gradient) $\partial J(\theta) / \partial \theta > 0$.
 $\theta^B - \alpha * \partial J(\theta) / \partial \theta < \theta^B$
 θ^B is moving to the left
 θ is decreasing

- α is called the learning rate or the step size
 - Too small
 - Take baby steps \rightarrow Take too long to converge
 - Too large
 - Can overshoot the minimum \rightarrow fail to converge

Simple Linear regression with GD

- Repeat until convergence :

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1)$$



- If we calculate the derivatives, the expression becomes:

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \quad \theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

- Convergence means either:
 - The cost function is no longer changing by more than ϵ .
 - The number of iterations is reached
- [Go further](#)

Ordinary Least Square Algorithm

- The Ordinary least square (OLS) approach chooses θ_0, θ_1 to minimize SSE

$$SSE = \sum_{i=1}^m (y^{(i)} - \theta_0 - \theta_1 * x^{(i)})^2$$

- In this method, we will minimize SSE by explicitly taking its derivatives with respect to the θ_0, θ_1 , and setting them to zero.
- The minimizing values can be shown to be:



$$\hat{\theta}_1 = \frac{\sum_{i=1}^m (x^{(i)} - \bar{x}) (y^{(i)} - \bar{y})}{\sum_{i=1}^m (x^{(i)} - \bar{x})^2}$$

$$\hat{\theta}_0 = \bar{y} - \hat{\theta}_1 * \bar{x}$$

- Where $\bar{y} = \frac{1}{m} \sum_{i=1}^m y^{(i)}$, $\bar{x} = \frac{1}{m} \sum_{i=1}^m x^{(i)}$ are the sample means

SSE is not perfect

- Which fit has larger SSE?



- Larger SSE doesn't necessarily mean worst fit → Need an other metric

R² Statistic

- R² Answers the questions: **“how much of variability in the output (y) is explained by the change in the input (x)”**
- To calculate R², we use the formula:

$$R^2 = \frac{TSS - SSE}{TSS} = 1 - \frac{SSE}{TSS}$$

$$TSS = \sum_{i=1}^m (y^{(i)} - \bar{y})^2$$

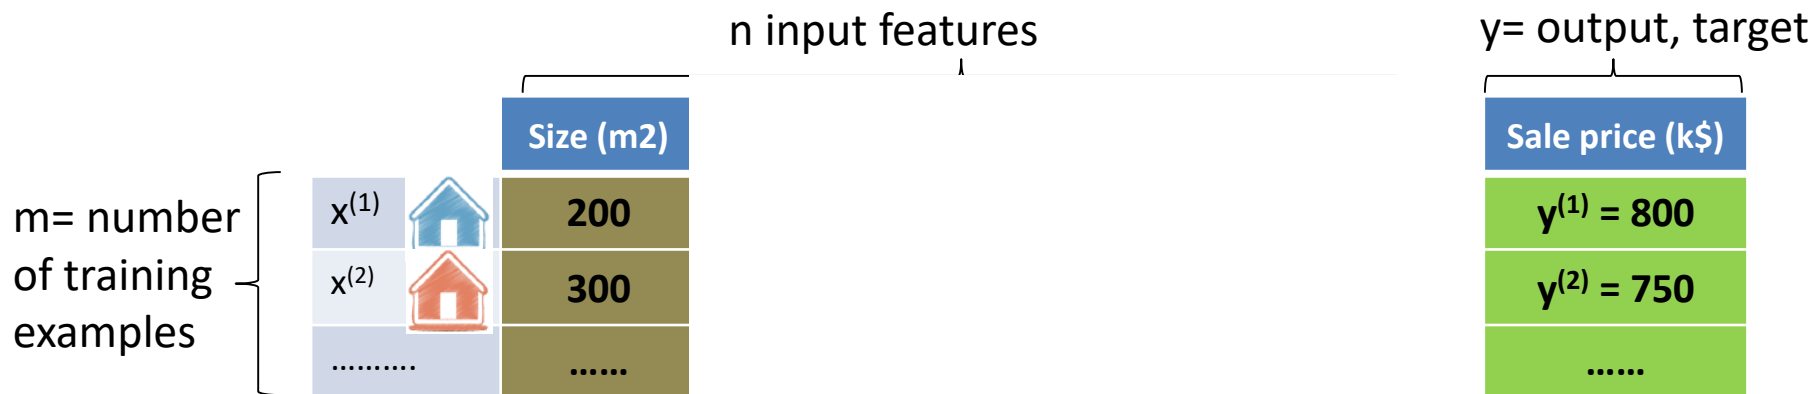
- An R² statistic that is close to 1 indicates that a large proportion of the variability in the response has been explained by the regression.
- A number near 0 indicates that the regression did not explain much of the variability in the response

1.2

Multivariate Linear Regression

Multiple Linear Regression

- In simple linear regression, we use one feature x to predict (y)
- In multiple linear regression, we have multiple features $X=(x_1, x_2, \dots, x_n)$



- $X^{(i)}$ is an n -dimensional feature vector
- $X^{(1)} = (200, 2010, 2, \dots)^T$ the feature vector of the first training example.
- $x_j^{(i)}$ is the value of feature j in the i^{th} training example. $x_2^{(1)} = 2010$

Multiple Linear Regression

- In simple linear regression, $h_{\theta}(x) = \theta_0 + \theta_1 x$
- In multiple linear regression, $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$

- For convenience of notation, define $x_0 = 1$ ($x_0^{(i)} = 1$)

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

- Rewrite in matrix notation

θ_0
θ_1
θ_2
...
θ_n

$$\theta, x \in R^{n+1}$$

x_0
x_1
x_2
...
x_n

$$h_{\theta}(x) = \sum_{i=1}^n \theta_i x_i = \theta^T x = \theta x^T$$

Multiple Linear Regression

- Cost function:

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Simple Linear Regression, n=1

$$\begin{aligned} &\text{Repeat } \{ \\ &\theta_0 := \theta_0 - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})}_{\frac{\partial}{\partial \theta_0} J(\theta)} \\ &\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)} \\ &\quad \text{(simultaneously update } \theta_0, \theta_1 \text{)} \} \end{aligned}$$

Multiple Linear Regression, n>1

$$\begin{aligned} &\text{Repeat } \{ \\ &\quad \theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \\ &\quad \text{(simultaneously update } \theta_j \text{ for } j = 0, \dots, n) \\ &\} \\ &\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)} \\ &\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)} \\ &\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_2^{(i)} \\ &\dots \end{aligned}$$

OLS for multiple features

- Cost function with matrix notations:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{2m} (X\theta - y)^T (X\theta - y)$$

- Where

$$X = \begin{bmatrix} \text{---} (x^{(1)})^T \text{---} \\ \text{---} (x^{(2)})^T \text{---} \\ \vdots \\ \text{---} (x^{(m)})^T \text{---} \end{bmatrix} \quad y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

- Derivatives with respect to θ $\nabla_{\theta} J(\theta) = X^T X \theta - X^T y$

- By setting them to zero $\theta = (X^T X)^{-1} X^T y$

When to use OLS or GD ?

- The following is a comparison of GD and the OLS (normal equation):


Gradient Descent	OLS (Normal Equation)
Need to choose alpha	No need to choose alpha
Needs feature scaling	No need for feature scaling
Needs many iterations	No need to iterate
$O(kn^2)$	$O(n^3)$, need to calculate inverse of $X^T X$
Works well when n is large	Slow if n is very large

- $n=10^4$ - 10^5 is usually the threshold of choosing GD over OLS
- Feature scaling helps converting the features to the same scale.
- For example, if x_i represents housing prices with a range of 100 to 2000 and a mean value of 1000, then,

$$x_i = \frac{price - 1000}{2000 - 100}$$

Polynomial Regression

- Polynomial regression is a particular case of multiple regression where the features are powers of one single feature x



- General model:

$$y = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \dots + \theta_p x^p + \varepsilon$$

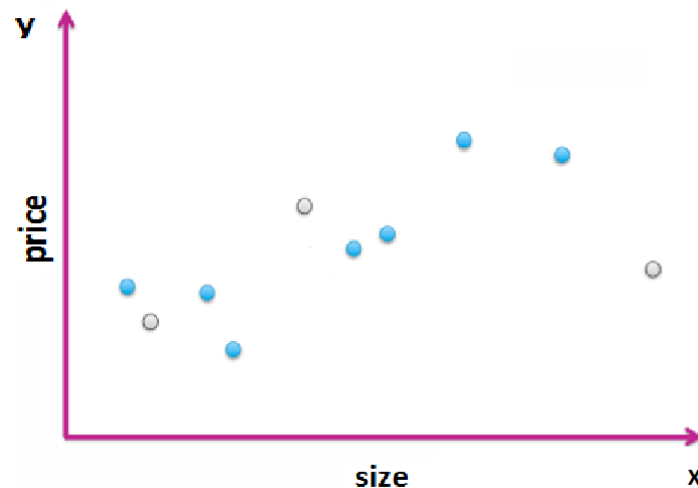
1.3

Assessing Performance

Assessing Performance

- This is about knowing how well the model will **generalize** to unseen data.
- One of the most used methods is splitting the data into train/test sets.


	Size	Price
Training set – 70%	2104	400
	1600	330
	2400	369
	1416	232
	3000	540
	1985	300
	1534	315
	1427	199
Test set 30%	1380	212
	1494	243



- Learn θ from training data (minimizing training error $J_{\text{train}}(\theta)$)
- Compute test error $J_{\text{test}}(\theta) = \frac{1}{2m_{\text{test}}} \sum_i^{m_{\text{test}}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$



Training error vs. model complexity

- Let's take polynomial regression as an example





Training error vs. model complexity

- Training error decreases with increasing model complexity




Testing error vs. model complexity

- Test error can be used as an approximation of the generalization error



Bias – Variance tradeoff


Model Complexity



- High bias model

The model does not capture enough
The structure of the training set
Parameters tend to be small

UNDERFIT




- High variance model

The model is too specific to the structure
Of the training set
Parameters tend to be very large

OVERFIT

Bias – Variance tradeoff



1.4

Regularization

Regularization

- It's about finding balance between:
 - How well the model fits the data
 - The magnitude of coefficients
- This is achieved by incorporating a penalty on weights θ in the cost function
- Ridge Regression (L_2 regularization)

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=0}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n (\theta_j)^2 \right]$$

- Lasso Regression (L_1 regularization)

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=0}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j| \right]$$

- λ is the regularization parameter:
 - Ridge: Encourages small weights θ but not exactly 0
 - Lasso: "Shrink" some weights θ exactly to 0

GD with Regularization

- Reminder of the L2 regularization cost function:

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

- Previously:
$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$
$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$
- With regularization:
$$\theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$
- α, λ are learning parameters to choose manually
- In practice: $(1 - \alpha\lambda/m)$ is between 0.99 and 0.95

Other Linear Regression Models

- Number of visiting customers to a website
- Product demand, inventory, failure, ...
- Stock pricing
- Insurance claims severity

“Remember that all models are wrong;
the practical question is how wrong do they
have to be to not be useful.”

George Box, 1987